

Adversarial Search

Dr. Steven Bethard

Computer and Information Sciences
University of Alabama at Birmingham

28 Jan 2016

Outline

1 Describing Games

- Games as Search
- Game Search Trees

2 Deterministic Games

- Minimax
- Alpha-Beta Pruning Example
- Approximate Solutions

3 Nondeterministic Games

- Describing Nondeterministic Games
- Expectiminimax
- Partially Observable Games

Outline

1 Describing Games

- Games as Search
- Game Search Trees

2 Deterministic Games

- Minimax
- Alpha-Beta Pruning Example
- Approximate Solutions

3 Nondeterministic Games

- Describing Nondeterministic Games
- Expectiminimax
- Partially Observable Games

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states

Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information		
Imperfect information		

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states

Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information	chess, checkers, go, reversi	backgammon, monopoly
Imperfect information	battleship, blind tic-tac-toe	bridge, poker, scrabble

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states

Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information	chess, checkers, go, reversi	backgammon, monopoly
Imperfect information	battleship, blind tic-tac-toe	bridge, poker, scrabble

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states

Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information	chess, checkers, go, reversi	backgammon, monopoly
Imperfect information	battleship, blind tic-tac-toe	bridge, poker, scrabble

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states

Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information	chess, checkers, go, reversi	backgammon, monopoly
Imperfect information	battleship, blind tic-tac-toe	bridge, poker, scrabble

Games as Search

Games as Search

Initial A board position and which player is to move

Actions Moves and resulting states


Goal A terminal state, where the game has ended

Score Based on utility function, e.g. +1 win, -1 lose

Game Types

	Deterministic	Stochastic
Perfect information	chess, checkers, go, reversi	backgammon, monopoly
Imperfect information	battleship, blind tic-tac-toe	bridge, poker, scrabble

Tic-Tac-Toe Game Tree



Outline

- 1 Describing Games
 - Games as Search
 - Game Search Trees
- 2 **Deterministic Games**
 - **Minimax**
 - **Alpha-Beta Pruning Example**
 - **Approximate Solutions**
- 3 Nondeterministic Games
 - Describing Nondeterministic Games
 - Expectiminimax
 - Partially Observable Games

Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player



Min Player


Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player




Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

Min Player



Minimax

Idea

- Expect other player to minimize your utility
- Choose action that maximizes the minimized utility

Max Player

3

Min Player


Minimax Exercise

Build Minimax Tree

- It is X's turn
- Scoring:
 - X Wins +1
 - X Loses -1
 - X and O Tie 0

O	X	X
X		
O	O	


Minimax Exercise



Minimax Properties

Max Player

Min Player



Complete?

Yes, if tree is finite

Optimal?

Yes, for optimal opponent

Time?

$O(b^m)$, all nodes in the tree


Space?

$O(bm)$, like depth first search

Minimax Properties

Max Player

Min Player



Complete? Yes, if tree is finite

Optimal? Yes, for optimal opponent


Time? $O(b^m)$, all nodes in the tree

Space? $O(bm)$, like depth first search

Minimax Properties

Max Player

Min Player



Complete? Yes, if tree is finite

Optimal? Yes, for optimal opponent


Time? $O(b^m)$, all nodes in the tree

Space? $O(bm)$, like depth first search

Minimax Properties

Max Player

Min Player



Complete? Yes, if tree is finite

Optimal? Yes, for optimal opponent


Time? $O(b^m)$, all nodes in the tree

Space? $O(bm)$, like depth first search

Minimax Properties


Max Player

Min Player




- | | |
|-----------|-----------------------------------|
| Complete? | Yes, if tree is finite |
| Optimal? | Yes, for optimal opponent |
| Time? | $O(b^m)$, all nodes in the tree |
| Space? | $O(bm)$, like depth first search |


Multiplayer Minimax




Multiplayer Minimax




Multiplayer Minimax




Multiplayer Minimax




Multiplayer Minimax



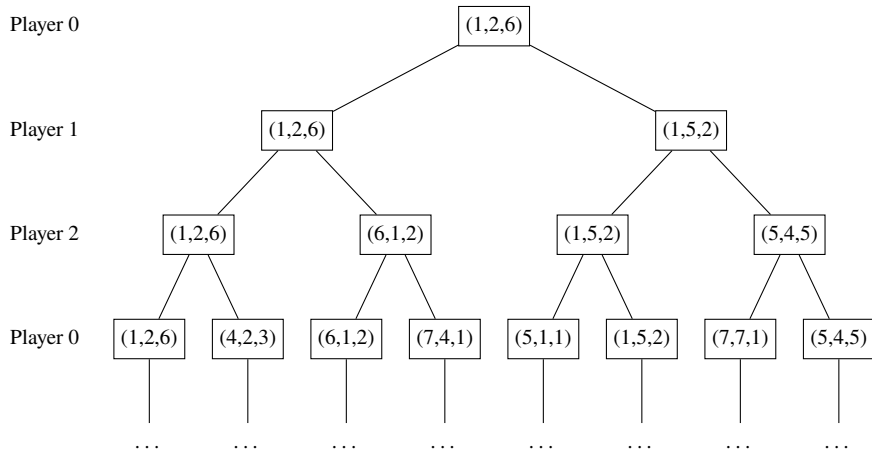
Multiplayer Minimax



Multiplayer Minimax



Multiplayer Minimax



Alpha-Beta Pruning

Idea

If m is a better choice, then n will never be reached

Alpha-Beta Bookkeeping

α Best value for Player (the highest value)

β Best value for Opp (the lowest value)


Player

Opponent

..
..
..

Player

Opponent



Alpha-Beta Pruning

Idea

If m is a better choice, then n will never be reached

Alpha-Beta Bookkeeping

α Best value for Player (the highest value)

β Best value for Opp (the lowest value)


Player

Opponent

..
..
..

Player

Opponent



Alpha-Beta Pruning

Max

$$\begin{array}{l} \geq -\infty \\ \leq +\infty \end{array}$$

Min

Max

Min


Alpha-Beta Pruning

Max


Min

Max


Min



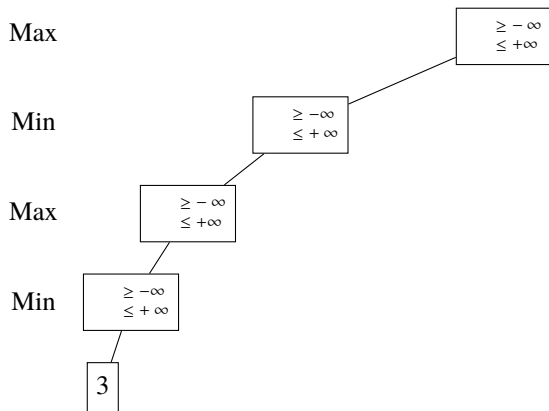
Alpha-Beta Pruning



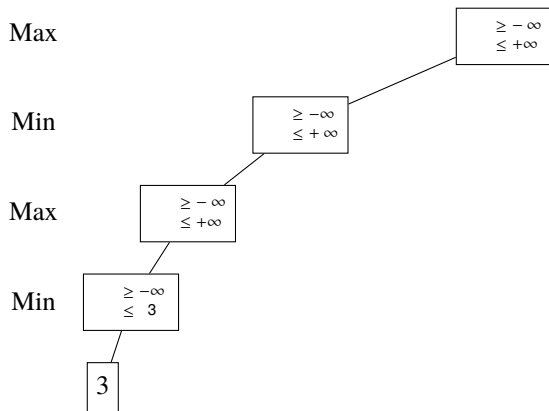
Alpha-Beta Pruning



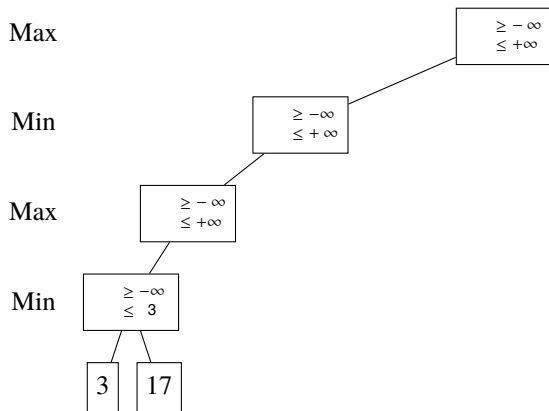
Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning



Alpha-Beta Pruning



Alpha-Beta Pruning




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max


Min

Max


Min




Alpha-Beta Pruning



Alpha-Beta Pruning



Alpha-Beta Pruning




Alpha-Beta Pruning

Max


Min

Max

Min



Alpha-Beta Pruning




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max


Min

Max


Min



Alpha-Beta Pruning



Alpha-Beta Pruning




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max

Min

Max

Min




Alpha-Beta Pruning

Max

Min

Max

Min



Alpha-Beta Pruning

Max


$$\boxed{3 \begin{array}{l} \geq 3 \\ \leq +\infty \end{array}}$$

Min

Max

Min


Alpha-Beta Pruning Properties



Time Complexity

- “Perfectly” ordered branches gives $O(b^{m/2})$
- Still exponential, but doubles solvable depth


Alpha-Beta Pruning Properties



Time Complexity

- “Perfectly” ordered branches gives $O(b^{m/2})$
- Still exponential, but doubles solvable depth

Alpha-Beta Pruning Properties



Time Complexity

- “Perfectly” ordered branches gives $O(b^{m/2})$
- Still exponential, but doubles solvable depth

Approximate Solutions

Problem

Even with Alpha-Beta Pruning, chess space is still 35^{50}

Solution

Don't search the entire tree:

```
if problem.is_terminal(state):  
    return problem.get_utility(state)
```

Example

- Given 100 seconds
- Given ability to explore 10^4 nodes/second
- Can explore to depth ≈ 8 ($10^6 \approx 35^{8/2}$)

Approximate Solutions

Problem

Even with Alpha-Beta Pruning, chess space is still 35^{50}

Solution

Don't search the entire tree:

```
if problem.is_terminal(state):  
    return problem.get_utility(state)
```

Example

- Given 100 seconds
- Given ability to explore 10^4 nodes/second
- Can explore to depth ≈ 8 ($10^6 \approx 35^{8/2}$)

Approximate Solutions

Problem

Even with Alpha-Beta Pruning, chess space is still 35^{50}

Solution

Don't search the entire tree:

```
if problem.is_past_cutoff(state):  
    return problem.get_estimated_utility(state)
```

Example

- Given 100 seconds
- Given ability to explore 10^4 nodes/second
- Can explore to depth ≈ 8 ($10^6 \approx 35^{8/2}$)

Approximate Solutions

Problem

Even with Alpha-Beta Pruning, chess space is still 35^{50}

Solution

Don't search the entire tree:

```
if problem.is_past_cutoff(state):  
    return problem.get_estimated_utility(state)
```

Example

- Given 100 seconds
- Given ability to explore 10^4 nodes/second
- Can explore to depth ≈ 8 ($10^6 \approx 35^{8/2}$)

Evaluation Functions

Necessary Properties

- Quickly computable
- For terminals, $\text{EVAL}(s)$ orders by utility
- For nonterminals, $\text{EVAL}(s)$ correlates with winning

Typical Approach

Weighted linear combination of features:

$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Evaluation Functions

Necessary Properties


- Quickly computable
- For terminals, $\text{EVAL}(s)$ orders by utility
- For nonterminals, $\text{EVAL}(s)$ correlates with winning

Typical Approach

Weighted linear combination of features:


$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Example Evaluation Function



Black to move

White slightly better



White to move

Black winning

Weighted Features

$$w_{\text{pawn}} \cdot f_{\text{pawn}}(s) = 1 \cdot (\text{white pawns} - \text{black pawns})$$

$$\dots = \dots$$

$$w_{\text{queen}} \cdot f_{\text{queen}}(s) = 9 \cdot (\text{white queens} - \text{black queens})$$

Design an Evaluation Function

Extended Tic-Tac-Toe

- $N \times N$ board
- K in a row wins

Must satisfy:

- Quickly computable
- $\forall x, t, o \in \text{TERMINALS},$
$$\text{UTIL}(x) = +1 \wedge \text{UTIL}(t) = 0 \wedge \text{UTIL}(o) = -1$$
$$\implies \text{EVAL}(x) > \text{EVAL}(t) > \text{EVAL}(o)$$
- $\forall s \in \text{NONTERMINALS}, \text{EVAL}(s)$ correlates with chance of +1

Deterministic Games in Practice

Checkers

- 1995: Chinook “Man-Machine World Champion” (1-0-31)
- 2007: Chinook’s creators “proved” it cannot lose
- End-game database for all ≤ 8 piece states

Chess

- 1997: Deep Blue defeated Garry Kasparov (2-1-3)
- Searches 6-40 plies; end-game database for ≤ 5 piece states
- 2002,’04,’06,’09,’11,’13: Junior wins World Computer Chess


Go

- Branching factor starts at 361 \Rightarrow Monte Carlo methods
- Computers still rank as advanced amateurs (6 dan)


Outline

- 1 Describing Games
 - Games as Search
 - Game Search Trees
- 2 Deterministic Games
 - Minimax
 - Alpha-Beta Pruning Example
 - Approximate Solutions
- 3 Nondeterministic Games
 - Describing Nondeterministic Games
 - Expectiminimax
 - Partially Observable Games

Nondeterministic Game: Backgammon



Backgammon Game Tree



Expectiminimax

Expected Values for Chance Nodes

$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$

Max



Chance

Min

Expectiminimax


Expected Values for Chance Nodes

$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$

Max

Chance


Min



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes


$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes

$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$



Expectiminimax

Expected Values for Chance Nodes

$$f(n) = \sum_{s \in \text{SUCCESSORS}(n)} P(s) \cdot f(s)$$

Max

3

Chance


Min

Expectiminimax Practice

A Simple Game

- Deck contains J, Q, K, K in a random order
- A card is drawn and Player 1 either:
 - Ends the game, or
 - Another card is drawn and Player 2 either:
 - Ends the game, or
 - Another card is drawn and the game ends
- Utility is determined by the last card drawn
- J scores -1, Q scores 0, and K scores +1

Expectiminimax Practice



Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood? [Node Likelihood](#)

Alpha-Beta Pruning? [Alpha-Beta Pruning](#)

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased

Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Using Expectiminimax

Expectiminimax Properties

Complete?	Yes, if tree is finite (both moves and “rolls”)
Optimal?	Yes
Time?	$O(b^m n^m)$, all nodes, all “roll” sequences

Consequences

Node Likelihood?	Decreases with depth
Alpha-Beta Pruning?	Effectiveness decreased



Real-World Example: TD-Gammon

- Depth 2-3 search, no Alpha-Beta pruning
- Neural network eval function trained by self-play

Utilities: Minimax vs. Expectiminimax

Max



Min



Utilities: Minimax vs. Expectiminimax

Max



Min



Max

Chc



Min



Utilities: Minimax vs. Expectiminimax

Max



Min



Max

Chc

Min



Evaluation function must be linear transform of true utility

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- Road B leads to a fork:
 - Take the left fork and you'll find a large heap of gold
 - Take the right fork and you'll receive your life

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- Road B leads to a fork:
 - Take the left fork and you'll find a large heap of gold
 - Take the right fork and you'll find an even larger one

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- Road B leads to a fork:
 - Take the left fork and you'll find a large heap of gold
 - Take the right fork and you'll be run over by a bus

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- **Road B** leads to a fork:
 - Take the **left fork** and you'll find a large heap of gold
 - Take the right fork and you'll be run over by a bus

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- Road B leads to a fork:
 - Take the left fork and you'll be run over by a bus
 - Take the right fork and you'll find a large heap of gold

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- **Road B** leads to a fork:
 - Take the left fork and you'll be run over by a bus
 - Take the **right fork** and you'll find a large heap of gold

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- Road A leads to a small heap of gold
- Road B leads to a fork:
 - Guess correctly and you'll find a large heap of gold
 - Guess incorrectly and you'll be run over by a bus

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns?

Problems with *averaging over clairvoyance*

- **Road A** leads to a small heap of gold
- Road B leads to a fork:
 - Guess correctly and you'll find a large heap of gold
 - Guess incorrectly and you'll be run over by a bus

Partially Observable Games

Missing knowledge on opponent state

- Poker, Blackjack, Bridge, etc.
- Just average over all possible unknowns? **No!**

Problems with *averaging over clairvoyance*

- **Road A** leads to a small heap of gold
- Road B leads to a fork:
 - Guess correctly and you'll find a large heap of gold
 - Guess incorrectly and you'll be run over by a bus

Partially Observable Games

Key Point

Value of an action is not the average across all states

Should be searching through a tree of belief states, and:

- Acting to obtain information
- Signaling to one's partner
- Acting randomly to minimize information disclosure

But in the Real World...

Most programs use Monte-Carlo estimation:

- Generate 100+ deals consistent with bidding
- Pick action that wins most on average

Partially Observable Games

Key Point

Value of an action is not the average across all states
Should be searching through a tree of belief states, and:

- Acting to obtain information
- Signaling to one's partner
- Acting randomly to minimize information disclosure

But in the Real World. . .

Most programs use Monte-Carlo estimation:

- Generate 100+ deals consistent with bidding
- Pick action that wins most on average

Partially Observable Games

Key Point

Value of an action is not the average across all states
Should be searching through a tree of belief states, and:

- Acting to obtain information
- Signaling to one's partner
- Acting randomly to minimize information disclosure

But in the Real World...

Most programs use Monte-Carlo estimation:

- Generate 100+ deals consistent with bidding
- Pick action that wins most on average

Key Points

Representing Games

- Multiple plies per round, one per player
- Stochastic games introduce chance nodes

Optimal Solutions

- (Expecti-)Minimax produces optimal actions
- Search belief states when information is incomplete

Approximate Solutions

- (Expecti-)Minimax explores the whole tree
- Approximations use utility estimates and cutoffs
- Chance dramatically reduces the depth explored